




# Flower Recognition: Dealing with Less Data via Few-Shot Learning

George Rahul  
Chopra Dhruv  
Agarwala Pratham





# /TABLE OF CONTENTS



**/01** MOTIVATION AND  
DATA AUGMENTATION

**/02** TRANSFER LEARNING


**/03** SIAMESE NETWORK

**/04** TRIPLET LOSS



/01


# MOTIVATION AND DATA AUGMENTATION



# THE FELINE FALLACY



# THE HUNGER OF NNs



# Can you predict with insufficient training data?

Human beings seem to be very good at recognising faces, cats, flowers etc. with only a few examples





# LITERATURE REVIEW


- Deep `transfer learning` (Tan et al., 2018).
- First `Siamese Network` for signature verification (Bromley et al., 1993).
- Face Verification using `Siamese Networks`, paper by `Facebook` (Taigman et al., 2014).
- `Triplet Loss` function, paper by `Google` (Schroff et al., 2015).



# <DATASET INFO!>

→

The Oxford Flowers 102 dataset is a consistent of 102 flower categories commonly occurring in the United Kingdom. Each class consists of between 40 and 258 images. The images have large scale, pose and light variations.




img\_shape: (500, 667, 3)




img\_shape: (500, 666, 3)



img\_shape: (670, 500, 3)



img\_shape: (500, 505, 3)



img\_shape: (500, 672, 3)





## <DATASET INFO!>



The dataset is divided into a training set, a validation set and a test set.

The `training` set and `validation` set each consist of `10 images per class` (totalling `1020` images each).

The `test` set consists of the remaining `6149` images (`minimum 20 per class`).



# DATA AUGMENTATION

We are performing following types of transformations

**RESIZING**

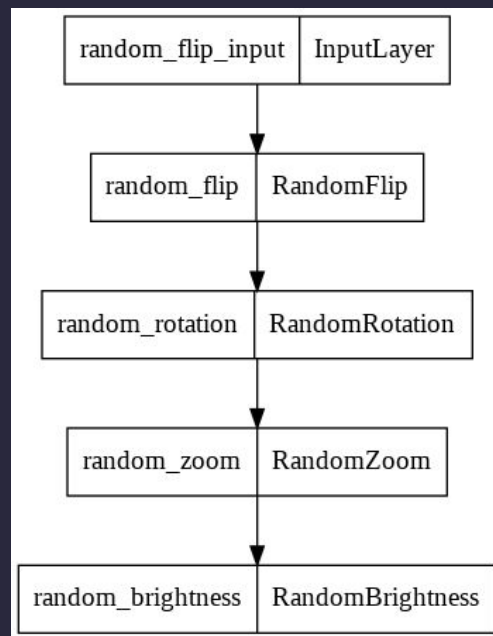
**RESCALING**

**FLIPPING**

**ROTATION**

**ZOOM/CROP**










**BRIGHTNESS**



# RESULTS AFTER AUGMENTATION

We used the `repeat` function of Tensorflow Dataset to repeat Our `Random Image Augmentation on the Dataset`.

Now we have over 5000 training images.



/02



# TRANSFER LEARNING

Using **Pre-trained** Networks  
and **fine-tune** to apply on  
our dataset



# TRANSFER LEARNING

- First step is add our `own classification layer` and train the network with keeping the `pretrained network freezed`.
- The second step is `fine-tuning` where we `unfreeze` a some layers and fit the model using `smaller learning rate( $1e^{-5}$ )`.



# DIFFERENT ARCHITECTURES

**/1** INCEPTION.V3


**/2** DENSENET 201

**/3** BigTransfer (BiT)

# LEARNING CURVE

We are using Adam optimizer, along with a Learning Rate Scheduler which decays the learning rate after a certain number of epochs.

The learning rate then decays exponentially for following epochs.



# TUNING MODEL ARCHITECTURE

1. Increasing Depth of the Network
2. Adding Batch Normalization Layers
3. Dropouts

We perform Bayesian Optimization Tuning using Keras tuner to obtain best parameters.





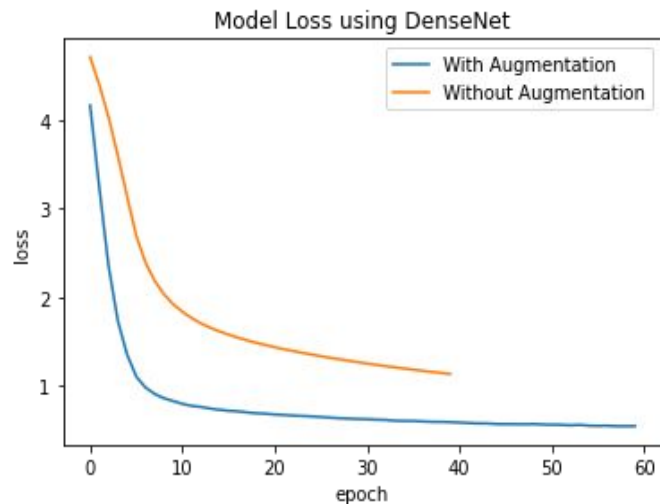
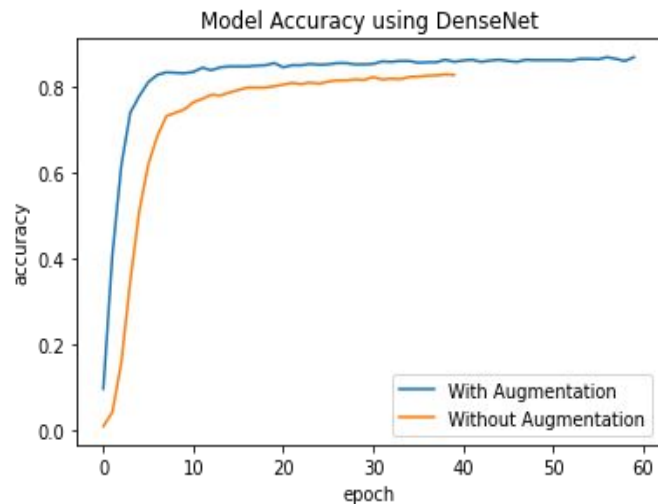
# /MODEL RESULTS ON TEST DATA

MODEL	ACCURACY	LOSS
Altered BiT	0.985	0.0817
BiT + Softmax	0.982	0.1011
Altered InceptionV3	0.75	1.03
DenseNet 201 + Softmax	0.90	0.4231



# EFFECT OF AUGMENTATION ON TRAINING


Below graphs represent validation loss across epochs trained for the DenseNet Model with and without Data Augmentation.





# /03

## FEW-SHOT LEARNING: SIAMESE NETWORK

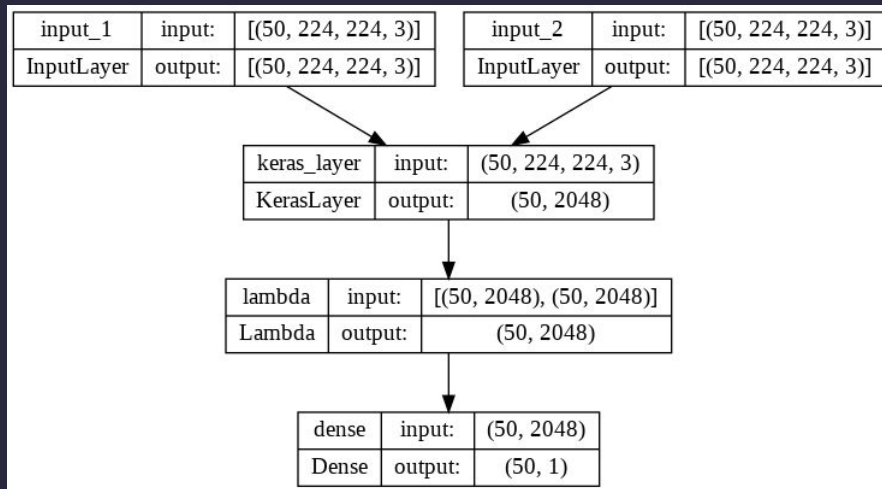


# ARCHITECTURE

We used the **BiG Transformer Feature Extractor** to encode the images.

And a **custom L2 layer** to calculate the distance between the encodings.

Network was trained with **binary cross entropy loss**.



# PREPARING THE DATA

We paired up the images **within the batches** to reduce complexity.

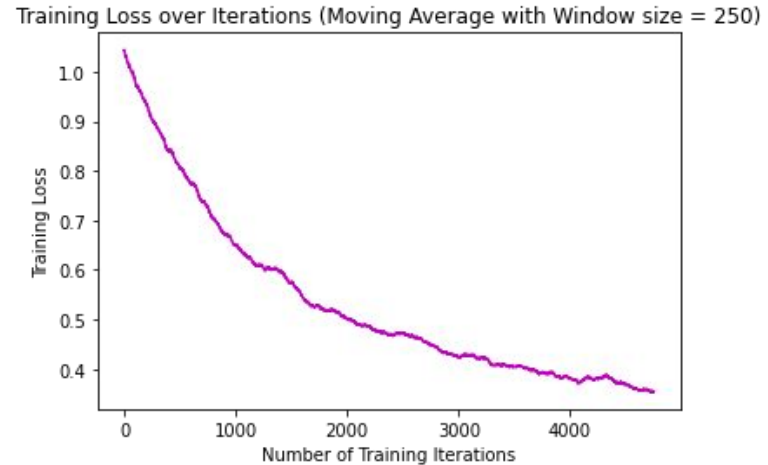
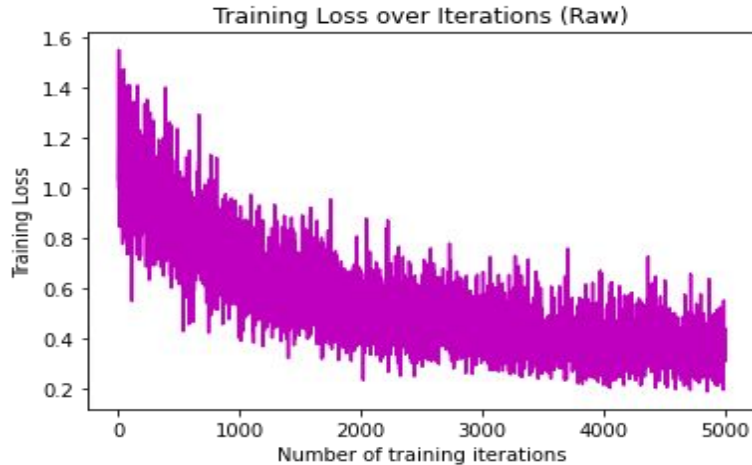
Each batch had a **1:1 ratio** of similar and dissimilar pairs.

We used batch size 50 and pairs were formed **by random sampling**.



# TRAINING THE MODEL

```
# Remember to compile the model!  
optimizer = SGD(learning_rate = 0.0001)  
model.compile(loss="binary_crossentropy", optimizer=optimizer)
```

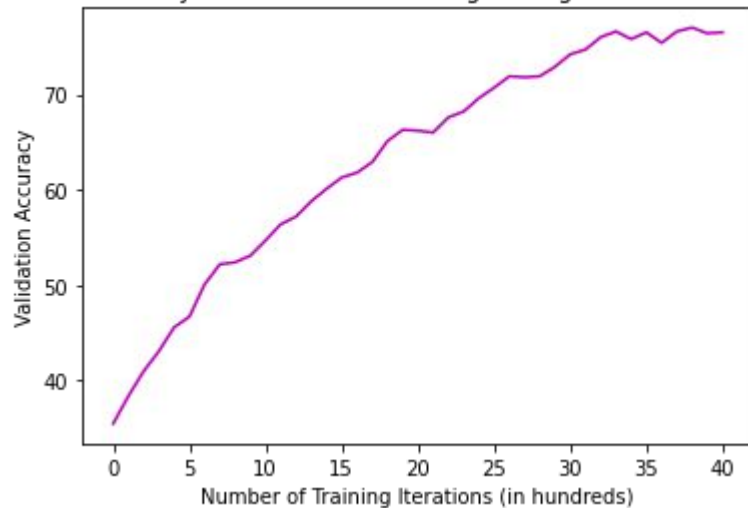


# VALIDATING THE MODEL

```
# N_way = N; n_val = k (num trials)
N_way = 5 # how many classes for testing one-shot tasks
n_val = 100 # how many one-shot tasks to validate on
```

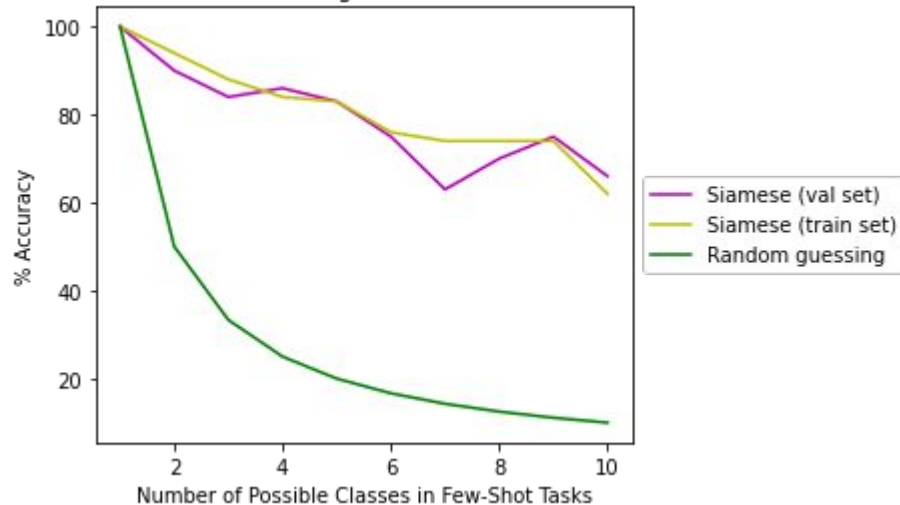


Validation Accuracy over Iterations (Moving Average with Window size = 10)



# VARYING THE NUMBER OF WAYS

Oxford Flowers Few-Shot Learning Performance of a Siamese Network





# /04

## FEW-SHOT LEARNING: TRIPLER LOSS



# /Deep Architecture


Augmented Image



Resnet

Bit 101x Pretrained Model

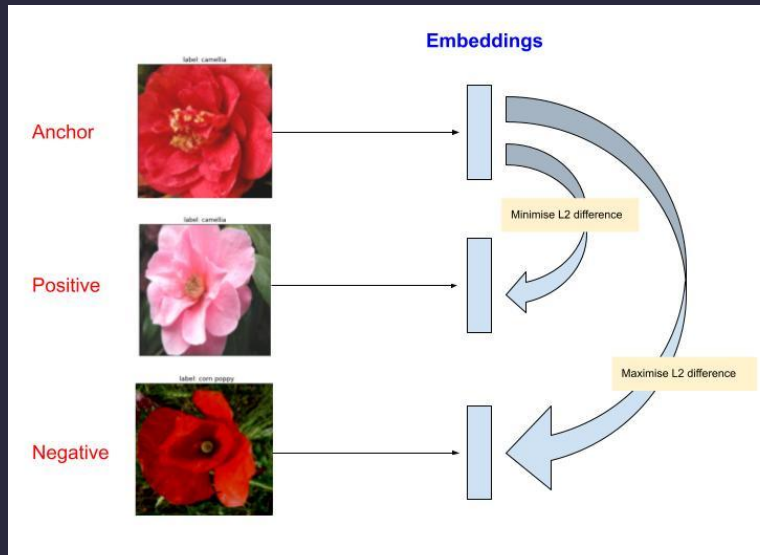
Dense Layers



Normalization Layer

Normalize

# /Some Terminologies



- **Anchor** Given Image
- **Positive** Category = Anchor Category
- **Negative** Category  $\neq$  Anchor Category

# /Improving the Model Loss

$$L2(rose_1, rose_2) < L2(rose_1, marigold_1)$$



Difference could be very small

$$L2(Anchor, Positive) < L2(Anchor, Negative) - margin$$




Ensure Loss is Positive

$$TripletLoss(A, P, N) = \max(0, L2(A, P) - L2(A, N) + margin)$$

# Sampling Triplets

& Semi-Hard Loss





# /MODEL RESULTS ON TEST DATA

	F1-Score	Recall	Precision
Non-weighted Avg	0.99	0.99	0.99
Weighted Avg	0.99	0.99	0.99

Achieved an accuracy of 0.9882907789884534





# /VISUALIZING THE EMBEDDINGS USING PCA

<https://projector.tensorflow.org>





# /DEPLOYED APPLICATION FOR DEMONSTRATION

<https://504910d1ba35053b.gradio.app/>







**THANK  
YOU**

